

Concurrency and Parallelism

- Concurrency and Parallelism are independent
- (true) parallelism is a hardware concept:
 - >1 processors; operations overlap in time
 - shared or local memory
- quasi-parallel:
 - single processor/memory
 - OS simulates parallelism via time-slicing
- concurrency is a software concept:
 - program objects are considered concurrent if they *could* be executed in parallel
 - concurrent operations *do not necessarily have to be executed in parallel*
 - many “real world” problems are easiest to program using a concurrent execution model
 - ◇ classic example: Operating system


Processes, Threads and Tasks

- a **process** is a program being executed; it contains.
 - program instructions
 - state (one program counter, one RTS, data, ...)
 - one *thread* of control per process
 - also called: **heavyweight process**, traditional process
- concurrent programs have possibly many active processes
→ they are described as *multithreaded*
- multithread programs have > 1 control thread
 - parallel processors – one process/CPU
 - ◇ OS manages synchronisation & context switching
 - single processor: many **threads** per process
 - ◇ each thread has own PC, stack, but share globals
 - ◇ controller *within the process* manages synchronisation & context switching
 - ◇ also called **lightweight process** or **task**



Standard control flow

- many units are *active*
- only one unit is executing at any time
- subroutines: master/slave relationship
- coroutines:
 - sharing of processing task
 - multiple entry points


Subprogram



Coroutine



Concurrent execution



- Main, A and B are independent
 - 3 tasks: Main, A, B
- Main waits for A & B to finish

Major concurrency issues

1. competition: e.g. ≥ 2 processes need to write database
2. cooperation: e.g. ≥ 2 processes share the processing task
3. *safety* and *liveness* of processes
 - (1, 2) require: *synchronisation* of actions of concurrent processes
 - (2) requires: *communication*: exchange of data between concurrent processes

Safety

Do all possible sequences result in the same state?

Consider two processes

$$P1: t := x; x := t+1 \quad (\equiv x:=x+1)$$
$$\text{P2: } u := x; x := u+2 \quad (\equiv x:=x+2)$$

If P1 & P2 run sequentially the final state is the same

$$P1 \text{ then } P2 \equiv P2 \text{ then } P1 \equiv x := x+3$$

Consider the cases if P1 & P2 run concurrently

Process 1	Process 2	
$t := x$	$u := x$	$\equiv x := x+1$
$x := t+1$	$x := u+2$	

Process 1	Process 2	
$t := x$	$u := x$	$\equiv x := x+2$
$x := t+1$	$x := u+2$	

Process 1	Process 2	
$t := x$		$\equiv x := x+3$
$x := t+1$	$u := x$	
	$x := u+2$	

Conclusion:

- *Interleaving* can lead to inconsistent state
- this is a competition synchronisation problem
 - both P1 and P2 compete to write to x

Liveness

Can a process complete in reasonable time?

Example: P1 and P2 both require exclusive access to X and Y

“Safe” interleaving where locking occurs in same order

Deadlock can occur if locking operations or in different order

Process 1	Process 2
lock X	
lock Y	lock X (<i>wait</i>)
update X,Y	
unlock Y	
unlock X	(<i>continue</i>)
	lock Y
	update X,Y
	unlock X
	unlock Y

Process 1	Process 2
lock X	
lock Y (<i>wait</i>)	lock Y
	lock X (<i>wait</i>)

- both examples illustrated competition concurrency
- the problem is one of *synchronisation*
- cooperation concurrency is similar but may also involve data exchange as well.
- data exchange: via shared data or messages
- a range of language features have been proposed to facilitate synchronisation and safe data exchange
 - *semaphores*: a synchronisation mechanism only
 - *monitors*: synchronisation + shared data
 - *message passing*: synchronisation + messages

Synchronisation via Semaphores

- proposed by Dijkstra (Algol68, Modula-2)
- two primitives: *P* (wait) and *V* (signal)
- a semaphore *s* is an integer variable
 - $\text{wait}(s) = \text{wait until } s > 0; s := s - 1$
 - $\text{signal}(s) = s := s + 1$
- implementation requires for each semaphore
 - FIFO queue of processes + integer variable
- usage: call wait before entering a *critical region* that requires exclusive access; call signal on exit from region.
- programmer is responsible for correct usage
- use of semaphores does not guarantee safeness of liveness; semaphores just provide a locking mechanism
- semaphores can be implemented in any language provided that wait and signal can be implemented as *atomic* operations.

Semaphore example

P1: wait(s); t := x; x := t+1; signal(s);

P2: wait(s); u := x; x := u+2; signal(s);

Main: var s:sem; s:=1; start(P1); start(P2)

Note: semaphore s specifies the number of processes that can be in the critical region. Normally set initially to 1 to guarantee exclusive access.

Process 1	Process 2	s
		1
wait(s)		0
t := x		0
	wait(s) (<i>wait</i>)	0
x := t+1		0
signal(s)		1
	(<i>resume</i>)	0
	u := x	0
	x := u+2	0
	signal(s)	1

Monitors

- Combine shared data and semaphore mechanism into a single ADT-like unit
- competition synchronisation is handled by allowing data access only via a monitor function.
 - no need for error prone semaphore
- cooperation synchronisation must be provided by semaphores
 - each language has different semaphore implementation
- Modula-2, Concurrent Pascal provides these facilities

Synchronisation via Message passing

- synchronisation is via a *rendezvous*
French: “meeting”
- rendezvous is a procedure call
- *client* calls the rendezvous
- *server* receives/accepts the rendezvous
- rendezvous can pass data as arguments
- ‘caller’ suspended during rendezvous

Ada example

```
procedure task_init is
  task type emitter is           -- declare "emitter" as a task
    entry init (c:character)     -- that will accept a rendezvous
  end emitter;                  -- called "init"

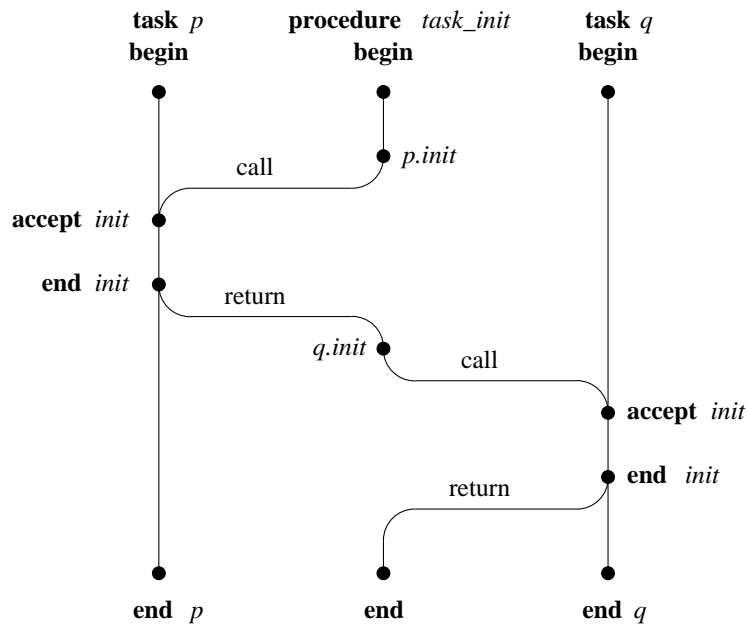
  p,q: emitter;                 -- create 2 tasks

  task body emitter is          -- emitter implementation
    me: character;
  begin
    accept init(c: character) do -- rendezvous
      me := c;                  -- critical region
    end init;
    put(me); new_line;
  end emitter;

begin
  p.init('p'); q.init('q'); put('r');
end task_init;
```

Produces output ‘p’ then ‘q’ then ‘r’

Control threads from example



Implementing Semaphores via message passing

```

task type binary_semaphore is
  entry wait;
  entry signal;
end binary_semaphore;

type sem_ptr is access binary_semaphore;

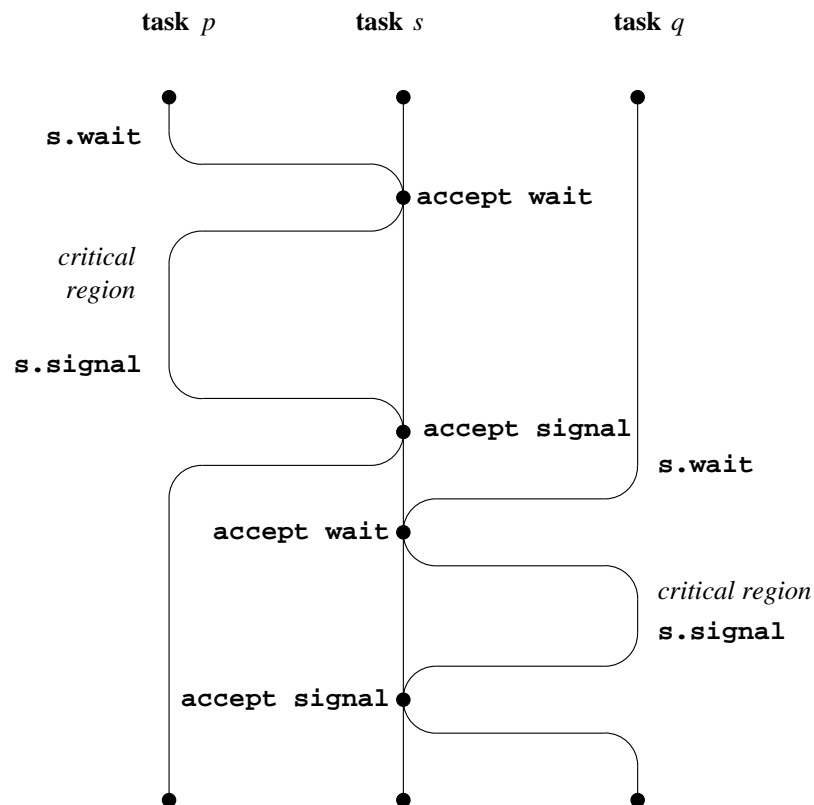
task body binary_semaphore is
begin
  loop
    accept wait;
    accept signal;
  end loop;
end binary_semaphore;

s : sem_ptr;

begin
  s := new binary_semaphore;
  ...
  s.wait;
  ...
  s.signal;
  ...
end;
  
```


Example

Tasks p and q use s to synchronise. p accesses shared data within critical region



Java threads

- Java objects can be concurrent if they either
 - inherit from predefined class `Thread`, or
 - implement the interface `Runnable`
- a threaded object defines two methods overriding those in `Thread`
 - `run` – which contains the body of the concurrent task, and
 - `start` – which calls `run`
- a task is started by calling `start` which immediately returns
- The Java runtime system contains a scheduler to control execution of runnable tasks.
- Threads can change priority (`getPriority/setPriority`)
- competition synchronisation is via `synchronized` methods or blocks
 - only one synchronised method/block can execute at a given time
 - thus private data in a class cannot have multiple readers/writers
 - similar functionality to a monitor

```
class Foo {  
    private int xyz[100];    // protected data goes here  
    public synchronized void writer (...) {...}  
    public synchronized int  reader (...) {...}  
}
```

- cooperation synchronisation uses `wait()` and `notify()`
 - `wait()` is called if thread cannot continue (perhaps because some other process is in a critical region); caller is suspended and placed in queue
 - `notify()` is called when a thread has completed a critical operation; other waiting tasks can then recheck and possibly continue.
- see Sebesta for examples.

Evaluation

- semaphores are primitive and error prone, but are sufficient
- monitors provide safer competition synchronisation
 - Java `synchronized` methods provide similar capability
- rendezvous does not need shared data
 - ⇒ can be used for true concurrency/distributed systems
- all cooperation synchronisation mechanisms are subject to programmer error: possible to generate deadlock situations

See Sebesta for further examples (e.g. Java, Ada95)