

G52CON: Concepts of Concurrency

Lecture 13 Message Passing

Brian Logan
School of Computer Science & IT
bsl@cs.nott.ac.uk

Outline of this lecture

- distributed processing implementations of concurrency
- message passing
 - one way vs two way communication
 - naming schemes
 - asynchronous vs synchronous communication
- client-server examples:
 - active monitors
 - self-scheduling disk driver

© Brian Logan 2002

G52CON Lecture 13: Message Passing

2

Implementations of concurrency

We can distinguish three types of implementations of concurrency:

- **multiprogramming**: execution of concurrent processes by timesharing them on a single processor;
- **multiprocessing**: the execution of concurrent processes by running them on separate processors which all access a shared memory; and
- **distributed processing**: the execution of concurrent processes by running them on separate processors which communicate by message passing.

© Brian Logan 2002

G52CON Lecture 13: Message Passing

3

Distributed processing

- *network architectures*, in which processors share only a communication network are becoming increasingly common: e.g., networks of workstations or multicomputers with distributed memory.
- the processes don't share a common address space, so they can't communicate via shared variables
- instead they communicate by *sending and receiving messages*

© Brian Logan 2002

G52CON Lecture 13: Message Passing

4

Message passing

Processes communicate by sending and receiving messages using special message passing primitives which include synchronisation:

- **send** *destination message*: sends *message* to another process *destination*
- **receive** *source message*: indicates that a process is ready to receive a message *message* from another process *source*

No special mechanisms are required for mutual exclusion, but new techniques are needed for condition synchronisation.

© Brian Logan 2002

G52CON Lecture 13: Message Passing

5

Approaches to message passing

Many different approaches to message passing have been proposed which differ in:

- whether communication is one way or two way;
- whether the naming of sources and destinations is direct or indirect
- whether the naming of sources and destinations is symmetrical or asymmetrical
- how **send** and **receive** operations are synchronised.

© Brian Logan 2002

G52CON Lecture 13: Message Passing

6

One way communication

If communication is *one way*

- process can only **send** or **receive** on a given channel in a single operation
- we need to use two channels to establish two way communication between processes

All the message passing primitives we will look at in this lecture use *one way communication*.

Naming sources and destinations

The naming of the source and destination of messages can be either:

- *direct*: using the names of processes; or
- *indirect*: using the names of channels; and

and

- *symmetrical*: both **send** and **receive** name processes or channels; or
- *asymmetrical*: only **send** names processes or channels and **receive** receives from any process or channel

Direct naming

In *direct naming*, unique names are given to all processes comprising a program

- in *symmetrical direct naming* both the sender and receiver name the corresponding process
- in *asymmetrical direct naming* the receiver can receive messages from any process

Indirect naming

Indirect naming uses intermediaries called channels or mailboxes

- in *symmetrical indirect naming* both the sender and receiver name the corresponding channel:
- in *asymmetrical indirect naming* the receiver can receive messages from any channel:

In this lecture, I will assume we are using *asymmetrical indirect naming*.

Processes and channels

Direct naming

symmetrical direct naming:
send *process message*
receive *process message*

asymmetrical direct naming:
send *process message*
receive *message*

Indirect naming

symmetrical indirect naming:
send *channel message*
receive *channel message*

asymmetrical indirect naming:
send *channel message*
receive *message*

Synchronising **send** and **receive**

- if a process tries to **receive** a message before one has been sent by another process, it will block until there is a message for it to read.
- the differences are mainly in the behaviour of the sending process:
 - asynchronous **send** : e.g., Unix sockets, `Java.net`
 - synchronous **send** : e.g., CSP, occam
 - remote invocation : e.g., extended rendezvous, RPC

Asynchronous Message Passing

If a process sends a message and continues executing without waiting for the message to be received, then the communication is termed *asynchronous*

- **send** operations are non-blocking;
- a sending process can get arbitrarily far ahead of a receiving process;
- message delivery is not guaranteed if failures can occur; and
- since channels can contain an unbounded number of messages messages have to be buffered.

Problems with asynchronous message passing

- the receiving process cannot know anything about the current state of the sending process
- the sending process has no way of knowing if the message was ever received unless the receiving process sends a reply
- it is hard to detect when failures have occurred
- buffer space is finite—if too many messages are sent either the program will crash, the buffer will overflow with loss of messages, or **send** operation will block

Synchronous message passing

If the sending process is delayed until the corresponding receive is executed, then the message passing is *synchronous*

- both the **send** and **receive** operations are blocking
- a process **sending** to a channel delays until another process is ready to receive from that channel;
- messages don't need to be buffered.

Problems with synchronous message passing

Synchronous message passing can result in reduced concurrency:

- whenever two processes communicate, at least one of them will have to block (whichever one tries to communicate first).
- concurrency is also reduced in some client-server interactions:
 - when a client is releasing a resource, there is usually no reason for it to wait until the server has received the release message
 - similarly, when a client writes to a device (e.g., a file or graphics display) it can usually continue without waiting for the server to receive the message.

Example: active monitors 1

```
// globally available names for n channels
channel request;
channel[] reply = new channel[n];

// Resource allocation process
process ResourceAllocator {
    list units = new list[MAXUNITS];
    queue pending;
    integer avail = MAXUNITS;
    integer clientID, unitID;
```

Active monitors 2

```
while (true) {
    receive request <clientID, kind, unitID>;
    if (kind == ACQUIRE) {
        if (avail > 0) {
            avail--;
            remove(units, unitID);
            send reply[clientID] <unitID>
        } else {
            insert(pending, clientID);
        }
    } else { // kind == RELEASE
        // free a unit of resource ...
```

Active monitors 3

```
// continued ...
} else { // kind == RELEASE
    if (empty(pending)) {
        avail++;
        insert(units, unitID);
    } else {
        remove(pending, clientID);
        send reply[clientID] <unitID>;
    }
}
}
```

© Brian Logan 2002

G52CON Lecture 13: Message Passing

20

Active monitors 4

```
// ith Client process of n ...
process Client {
    integer unitID;
    send request <i, ACQUIRE, null>;
    receive reply[i] <unitID>;
    // use the resource unitID, then release it ...
    send request <i, RELEASE, unitID>;
}
```

© Brian Logan 2002

G52CON Lecture 13: Message Passing

21

Active monitors 5

This is one way to program a simple monitor as an active process rather than a passive collection of procedures:

- we get mutual exclusion because only the server process can access its own local variables
- the monitor procedures typically get turned into the branches of an `if` or `switch` statement, so only one 'monitor procedure' can be active at a time, and the monitor procedures run with mutual exclusion.
- condition synchronisation is programmed with normal variables—conditions are re-evaluated on the arrival of a new message.

© Brian Logan 2002


G52CON Lecture 13: Message Passing

22

Example: scheduling disk access

- data is stored in *blocks* which are arranged in concentric *tracks*
- tracks in the same relative position on different platters form a *cylinder*
- the address of a block consists of a cylinder, track and offset from a start point
- data is accessed by positioning a read/write head over the appropriate track and waiting until the required block passes under the head.

A moving head disk



© Brian Logan 2002

G52CON Lecture 13: Message Passing

23

Disk access time

- disk access time depends on:
 - the seek time to move the heads to the appropriate cylinder
 - the rotational delay waiting for the block to pass under the heads
 - the data transmission time
- the time required to move the heads depends on the distance between the two cylinders—seeking even one cylinder takes much longer than the maximum rotational delay
- to reduce the *average* disk access time, we need to minimise head motion and hence seek time.

© Brian Logan 2002

G52CON Lecture 13: Message Passing

24

Scheduling disk requests

The key idea is to service the disk requests *out of order* when two or more clients want to access the disk at about the same time.

Suppose we have a sequence of requests for cylinders:

100, 1, 101, 2, ...

with the heads initially at cylinder 100. If we process these in the order of arrival, the heads travel a total of 300 cylinders.

If we process them in the order:

100, 101, 2, 1 ...


the heads travel a total of 101 cylinders. Furthermore the client requesting cylinder 1 only has to wait for the heads to move an extra two cylinders.

© Brian Logan 2002

G52CON Lecture 13: Message Passing

25

Separate scheduler




© Brian Logan 2002

G52CON Lecture 13: Message Passing

26

Scheduler as intermediary




© Brian Logan 2002

G52CON Lecture 13: Message Passing

27

Self-scheduling disk driver



© Brian Logan 2002

G52CON Lecture 13: Message Passing

28

Message passing solution 1

```

// shared channels
channel request;
channel[] reply = new channel[n];

// ith Client process of n ...
process Clienti {
    disk_block dblock;
    integer track, offset;
    // request a block of data from the disk
    send request <i, track, offset>;
    receive reply[i] <dblock>;
    // use the data from the disk ...
}
  
```

© Brian Logan 2002

G52CON Lecture 13: Message Passing

29

Message passing solution 2

```

// Self-Scheduling Disk Driver process
process DiskDriver {
    queue saved;
    disk_block dblock;
    integer clientID, track, offset, nsaved = 0;
    while(true) {
        while(!empty(request) or nsaved == 0) {
            // wait for first request or receive another
            receive request <clientID, track, offset>;
            insert(saved, <clientID, track, offset>);
            nsaved++;
        }
        // select the pending request closest to the ...
    }
}
  
```

© Brian Logan 2002

G52CON Lecture 13: Message Passing

30

Message passing solution 3

```

// select the saved request closest to the
// current head position and access the disk
remove(saved, <clientID, track, offset>);
nsaved--;
dblock = read(track, offset);
send reply[clientID] <dblock>;
}
}
  
```

© Brian Logan 2002

G52CON Lecture 13: Message Passing

31

Summary

- on a shared memory machine, procedure calls and operations on condition variables are more efficient than message passing primitives
- most distributed systems are based on message passing since it is more natural and more efficient than simulating shared memory on a distributed memory machine
- neither asynchronous nor synchronous message passing have yet found their way into a widely accepted general purpose programming language
- message passing in concurrent programs remains at the level of OS or library calls.

The next lecture

Remote invocation

Suggested reading:

- Andrews (2000), chapter 8;
- Ben-Ari (1982), chapter 6;
- Burns & Davies (1993), chapter 5;
- Andrews (1991), chapters 9.